

Appendix for PinpointQA

Due to space constraints in the main paper, this appendix supplies supplementary details and analyses. We begin with additional information on dataset construction and a more in-depth description of the evaluation metrics. We then examine representative failure cases and provide implementation details for the human localization study with FSD guidance. The appendix concludes with a brief discussion of the benchmark’s scope, limitations, and intended use.

A. Details of Dataset Construction

1. Scene Curation and Target Vocabulary
2. Reference Objects and Supporting Surface
3. QA Generation
4. Quality Control

B. Details of Evaluation Metrics

1. FSD Evaluation Setting
2. SSP Evaluation Setting

C. Representative Qualitative Examples

1. TPV Failure Case
2. NRI Failure Case
3. FSD Failure Case
4. SSP Failure Case

D. Details of Human Localization with FSD Guidance

1. Interface and Protocol
2. Scoring Scheme

E. Discussion

1. Intended Use
2. Limitations

A Details of Dataset Construction

This section supplements the construction pipeline described in the main paper by providing the complete target vocabulary and concise details about how the intermediate spatial representations are used across the four tasks.

A.1 Scene Curation and Target Vocabulary

Scene curation first identifies eligible target objects using a pre-defined category vocabulary. The vocabulary contains 109 source label variants describing portable or movable indoor objects, particularly objects that can be held in hand or are commonly placed on tables. After converting each source label to lowercase, we retain it only when it exactly matches one of these variants.

For benchmark statistics, the 109 source label variants are counted as 102 target categories. We group only source labels that denote the same reported category. Specifically, *phone*, *smartphone*, and *mobile phone* are grouped as *phone*; *remote control*, *tv remote*, *remote controller*, *remote*, and *tv controller* are grouped as *remote control*; and *mouse pad* and *mousepad* are grouped as *mousepad*. These three groups merge ten source label variants into three reported categories. Each of the remaining 99 source label variants corresponds to a separate category, yielding 102 categories in total. This grouping is used only when reporting category statistics, while the

original source labels are retained in the generated questions and answers.

Equivalent Labels Used for Category Statistics

phone, smartphone, mobile phone → phone
remote control, tv remote, remote controller, remote, tv controller → remote control
mouse pad, mousepad → mousepad

Source Label Vocabulary Used for Exact Matching

bag, backpack, book, bottle, bowl, candle, cup, dumbbell, hair dryer, hat, headphones, keyboard, laptop, mouse, purse, shoe, tissue box, water bottle, alarm clock, ball, stuffed animal, plate, clock, cd case, coffee kettle, paper towel roll, scale, soap dish, soap dispenser, telephone, toaster, tray, water pitcher, fire extinguisher, paper cutter, plunger, dustpan, broom, power strip, toilet paper, guitar, toaster oven, paper bag, key, phone, smartphone, mobile phone, earbuds, earphones, wireless headphones, headset, headphone case, remote control, tv remote, remote controller, remote, tv controller, game controller, charger, power adapter, laptop charger, phone charger, power bank, wireless charger, power plug, adapter, dongle, usb hub, extension cord, power cord, cable, wallet, handbag, tote bag, glasses, eyeglasses, glasses case, mouse pad, mousepad, pen, pencil, marker, highlighter, notebook, notepad, mug, spray bottle, lotion bottle, deodorant bottle, toothbrush, razor, comb, hairbrush, scissors, stapler, tape dispenser, glue bottle, calculator, lanyard, pen holder, pencil holder, pencil case, post it note, sticky note, phone stand, tissue, eraser, soap, power extension

A.2 Reference Objects and Supporting Surface

Within each intermediate spatial representation, supporting surfaces and reference objects serve different roles. The supporting surface is the primary spatial anchor selected for the target when a suitable candidate is available, whereas reference objects provide additional localization cues through their identities, spatial relations, and distances to the target. NRI excludes the selected supporting surface and asks for the nearest remaining valid reference. FSD uses the supporting surface as its main anchor when available and supplements it with nearby reference objects. SSP preserves the supporting surface and references as separate fields in the structured output. These distinct roles allow the same intermediate spatial representation to support the different localization requirements of the four tasks.

A.3 QA Generation

The four task formats are instantiated independently from the same intermediate spatial representations. For TPV, positive questions use retained target labels, while negative questions use vocabulary labels that are absent from the corresponding scene. For NRI, the correct answer is the closest valid reference after excluding the supporting surface, floor and wall categories, and objects in direct contact with the target. Each NRI question contains three distinct distractors, and cases with a near tie between the closest reference labels are removed. For FSD, the selected supporting surface is used as the main anchor when available; otherwise, the closest valid reference is used. Up to two additional references are included with their spatial relations and distances rounded to the nearest centimeter. For SSP, the target, the selected supporting surface, and up to two closest references are encoded in a parsable JSON object, with distances retained to one decimal place. The progression across TPV, NRI, FSD, and SSP therefore reflects increasingly demanding spatial and output requirements, while each task is generated directly from the shared intermediate spatial representation.

A.4 Quality Control

Quality control combines automatic validity checks with iterative manual review. Automatic checks remove invalid or structural labels, ambiguous target instances, and QA pairs that do not satisfy the required output constraints. TPV answers must contain a valid *Yes* or *No* label. NRI questions must contain four distinct options and a uniquely determined nearest valid reference. FSD answers must provide a complete natural language description grounded in at least one valid spatial anchor. SSP answers must be directly parsable as JSON objects containing the required `target`, `support_surface`, and `references` fields. Samples or scenes that fail these checks are excluded.

In each manual review round, we inspect approximately 100 QA pairs from 10–15 scenes. The review examines whether NRI excludes the supporting surface from its reference candidates, whether FSD preserves the intended supporting surface and nearby references, whether distances and semantic relations remain consistent with the underlying geometric measurements, whether the question and answer templates are natural and well formed, and whether the resulting localization is plausible to a human observer. Problems identified during a review round are used to revise the QA generation rules, filtering criteria, and templates before the next round, while the underlying intermediate spatial representations remain unchanged.

B Details of Evaluation Metrics

B.1 FSD Evaluation Setting

FSD predictions are evaluated by GPT-5.4 using a fixed scoring prompt and a temperature of zero. For each sample, the judge receives the target label, the ground truth description, and the model prediction. It evaluates whether the prediction preserves the supporting surface, reference objects, spatial relations, and numeric distances required to localize the target. The judge is instructed to assess only information supported by the ground truth and not to

reward additional references merely because they are plausible in the scene.

The score contains five dimensions with a maximum total of 10 points. Main location and supporting surface receive up to 3 points, reference objects receive up to 2 points, spatial relations receive up to 2 points, numeric distances receive up to 2 points, and clarity receives 1 point. The scoring criteria for these dimensions are summarized below.

FSD Scoring Rubric

Main location and supporting surface, 0–3 points. Three points are assigned when the main location and supporting surface are correct and the overall placement agrees with the ground truth. Two points are assigned when the supporting surface or primary anchor is correct but important localization details are missing. One point is assigned when the answer identifies only a broad related region without a reliable supporting surface or precise placement. A clearly incorrect supporting surface or scene region receives zero points.

Reference objects, 0–2 points. Two points are assigned when at least two correct and relevant references are used for localization. One point is assigned when exactly one correct reference is used meaningfully. Unsupported, incorrect, vague, or merely mentioned references receive no credit.

Spatial relations, 0–2 points. Two points are assigned when the important relations are correct. One point is assigned when at least one important directional relation is correct but another is incomplete, less precise, or incorrect. Predictions in which the important relations are missing, reversed, or replaced mainly by vague proximity receive no credit. In particular, replacing a directional relation in the ground truth with *near* does not constitute an equivalent relation.

Numeric distances, 0–2 points. Two points are assigned when the answer contains usable numeric distances that are tied to correct references and are roughly consistent with the ground truth. One point is assigned when at least one usable distance is correctly grounded but remains incomplete or moderately inaccurate. A distance receives no credit when it has no usable unit, is not tied to a correct reference, or is replaced by a vague expression. The judge uses approximate tolerances of 5 cm for ground truth distances no greater than 10 cm, 10 cm for distances between 10 cm and 30 cm, and 15 cm for distances greater than 30 cm.

Clarity, 0–1 point. One point is assigned when the answer is clear and readable. This dimension evaluates presentation only and cannot compensate for incorrect spatial content.

The judge is required to return a JSON object containing the total score, the five dimension scores, a concise explanation, and a set of error tags. The allowed tags are listed below:

```
wrong_main_location, wrong_support_surface,
missing_key_reference, wrong_spatial_relation,
missing_numeric_distance, wrong_numeric_distance,
unclear_expression, and hallucinated_reference.
```

If both the main location score and the reference object score are zero, the total score cannot exceed 2. Predictions containing almost no usable localization information are subject to the same upper bound. After the response is returned, the evaluator checks whether the total score equals the sum of the five dimension scores and replaces an inconsistent total with the computed sum. The corrected total is divided by 10 to obtain the normalized FSD score.

B.2 SSP Evaluation Setting

SSP is evaluated programmatically because its output has a pre-defined JSON structure. The evaluator first attempts to parse the complete model output as a JSON object. If this fails, it removes surrounding code fences and attempts to recover a JSON object from any additional text. A prediction receives a score of zero when no JSON object can be recovered. The recovered object is checked for a `support_surface` field and a `references` list. The alternative field name `primary_surface` is also interpreted as `support_surface`. Missing or malformed fields receive no credit for their corresponding components.

Before comparison, object and supporting surface labels are converted to lowercase, underscores and hyphens are converted to spaces, unnecessary quotation marks and parenthetical text are removed, and repeated spaces are collapsed. A conservative singularization rule and a fixed synonym map are then applied to common equivalent forms. For example, *mobile phone*, *smartphone*, and *cell phone* are normalized to *phone*; *night stand* is normalized to *nightstand*; and *table top* is normalized to *table*. Relation tokens undergo similar normalization so that forms such as *beside* and *adjacent to* map to *next_to*, while *close to* and *nearby* map to *near*. These operations handle equivalent surface forms without permitting unrestricted semantic matching.

Since the queried target is already specified in the question, the target field is not included in the SSP score. The evaluator scores the supporting surface, reference object identity, spatial relation, and distance. Supporting surfaces must match exactly after normalization. Predicted references are aligned one to one with the ground truth reference slots by selecting the assignment with the highest agreement in object identity, relation, and distance. Unmatched ground truth slots receive zero credit.

A reference object receives full credit only when its normalized label matches the corresponding ground truth label. Relation credit is computed only after the reference object has been matched. An exact relation match receives full credit, while the pairs *next_to* and *near*, *under* and *below*, and *on* and *attached_to* receive half credit. All other relation mismatches receive zero credit.

Distance credit requires both a correct reference object and positive relation credit. Let e denote the absolute difference between the predicted and ground truth distances. The exact distance thresholds are:

Ground truth distance	Full credit	Half credit
$d \leq 10$ cm	$e \leq 5$ cm	$e \leq 10$ cm
$10 < d \leq 30$ cm	$e \leq 10$ cm	$e \leq 15$ cm
$d > 30$ cm	$e \leq 15$ cm	$e \leq 20$ cm

Errors beyond the half credit threshold receive zero distance credit. For each reference slot, the distance credit is multiplied by its relation credit. Consequently, a half correct relation can

contribute at most half of the available distance credit, even when the numeric distance itself falls within the full credit threshold. The object, relation, and distance components are averaged over the ground truth reference slots.

The final normalized SSP score is computed as

$$S_{SSP} = 0.4S_{surface} + 0.2S_{object} + 0.2S_{relation} + 0.2S_{distance}.$$

This procedure separates structural validity from spatial correctness: a model may produce a parsable JSON object while still receiving a low score because its supporting surface, references, relations, or distances are incorrectly grounded.

C Representative Qualitative Examples

This section examines representative errors from three scenes with different target objects: *mouse* in scene 062e5a23a6, *mobile phone* in scene 0b031f3119, and *headphone case* in scene 281ba69af1. Figures 1, 2, and 3 show that errors can occur in different localization components, including target grounding, nearest reference identification, supporting surface prediction, reference grounding, spatial relations, numeric distances, and consistency across structured fields. Green text indicates content that agrees with the ground truth, while red text indicates incorrect or unsupported content.

C.1 TPV Failure Case

Figure 1 shows a target grounding error for *mouse* in scene 062e5a23a6. The ground truth confirms that the mouse appears in the scene, whereas Cambrian-S-7B predicts *No*. Because TPV requires neither reference selection nor structured output, this error can be attributed directly to incorrect target presence verification rather than to a later formatting or relation error. The example therefore illustrates that a model may process the broader scene while still failing to identify the queried small object.

C.2 NRI Failure Case

The same scene also provides a representative NRI error. As shown in Figure 1, the correct answer for the nearest valid reference to the mouse is *office chair*, but GPT-5.4 selects *laptop bag*. The supporting surface is excluded from the candidate set in this task, so the error specifically reflects an incorrect comparison among the remaining reference candidates. Although the predicted object is present in the scene and appears plausible as a reference, it is not the closest valid reference defined by the ground truth.

C.3 FSD Failure Case

The FSD examples show that a fluent description can still contain incorrect localization information. In Figure 1, LLaVA-OneVision-1.5-8B describes the *mouse* using a laptop and keyboard, while the ground truth uses the office chair and glasses case as the relevant references. The predicted relations and distances are therefore attached to objects that do not match the ground truth.

Figure 2 shows a similar error for *mobile phone*. Kimi K2.5 describes the target relative to a monitor, power outlets, and keyboard, omitting the tripod that provides the closest ground truth reference. InternVL3.5-8B-Instruct retains the table label but replaces the remaining reference configuration with a keyboard and mouse,

together with distances that differ from the ground truth. In Figure 3, GPT-5.4 places the *headphone case* on a desk but replaces the speaker and monitor with a cloth, tissue packet, and plant. LLaVA-OneVision-1.5-8B instead places the target on the floor and introduces a chair and towel that are not part of the ground truth localization.

Across these cases, the generated answers remain readable, but their supporting surfaces, reference objects, relations, or distances are incorrect. The errors therefore arise from failure to preserve the grounded spatial information required by FSD rather than from a lack of language fluency.

C.4 SSP Failure Case

The SSP examples demonstrate that producing a parsable JSON object does not guarantee correct structured spatial prediction. In Figure 1, SenseNova-SI-1.3-InternVL3-8B predicts *chair* as the supporting surface instead of *table* and replaces the ground truth references, *office chair* and *glasses case*, with *chair* and *desk*. The same chair is used as both the supporting surface and a reference object, creating a conflict between the structured fields. Its predicted relations and distances also differ from the ground truth.

Figure 2 presents two additional structured prediction errors. Kimi K2.5 produces valid JSON but changes the supporting surface from *table* to *desk* and replaces the ground truth references with *keyboard* and *monitor*. InternVL3.5-8B-Instruct preserves the table as the supporting surface but predicts *keyboard* and *mouse* as references, with incorrect relations and distances. Similarly, in Figure 3, GPT-5.4 predicts *tissue box* and *door* instead of *speaker* and *monitor*, while SenseNova-SI-1.3-InternVL3-8B changes the supporting surface to *chair* and introduces a different reference configuration.

These cases cover supporting surface mismatches, reference hallucination or drift, relation inconsistencies, inaccurate distances, and conflicts across structured fields. They also show why low SSP performance cannot be explained only by failures to generate valid JSON: an output may satisfy the requested structure while its localization fields remain inconsistent with the ground truth.

D Details of Human Localization with FSD Guidance

D.1 Interface and Protocol

We implement the human localization study described in Section 4.4 of the main paper using a web interface. For each sample, the interface presents the question together with the sampled video frames. Participants browse the frames, select the frame in which they believe the target appears, click its estimated location, and save their response. The visual interaction remains the same across the three settings, while the available textual guidance differs. The Unguided setting contains no spatial description. The Model-Assisted setting displays the FSD generated by Qwen3-VL-8B-Instruct-SFT, whereas the GT-Assisted setting displays the corresponding ground truth FSD. Figure 4 shows the interface in the GT-Assisted setting.

Timing begins when the participant starts the trial and the video frames are revealed, and ends when the participant saves the response. The recorded completion time therefore includes both frame selection and point localization. Before the study, we annotate the

target location using the same interface in annotation mode. For each sample, this process records the reference frame index and the target coordinates normalized by the image dimensions.

D.2 Scoring Scheme

Each submitted response consists of a selected frame index and normalized click coordinates (x, y) . Let (x^*, y^*) denote the annotated target coordinates. When the selected frame is correct, the normalized Euclidean distance between the submitted and annotated points is computed as

$$d = \sqrt{(x - x^*)^2 + (y - y^*)^2}.$$

An incorrect frame receives a score of 0. For a correct frame, we use a fixed radius of $r = 0.12$ and compute the score as

$$s = \max\left(0, 1 - \left(\frac{d}{r}\right)^2\right).$$

The score decreases quadratically with the distance from the annotated target and becomes zero when the distance reaches r .

For N evaluated responses, the reported localization accuracy is

$$\text{Accuracy} = \frac{100}{N} \sum_{i=1}^N s_i.$$

The reported completion time is the mean elapsed time across the same responses. This protocol jointly evaluates correct frame selection and accurate target localization. As stated in the main paper, the study provides auxiliary evidence about the usefulness of FSD guidance for human localization rather than a deployment evaluation.

E Discussion

E.1 Intended Use

PinpointQA is intended as a focused diagnostic benchmark and supervision resource for small object-centric spatial understanding in indoor videos. It does not aim to cover general video understanding or provide a corpus for large scale web pretraining. Instead, its four tasks examine a progressive capability chain from target presence verification to reference grounding, fine-grained spatial description, and structured spatial prediction. This design supports analysis of whether a model can identify a small target and accurately recover its supporting surface, nearby reference objects, spatial relations, and distances from video observations.

In addition to evaluation, the benchmark can provide supervision for the spatial QA tasks defined in this work. The improvements obtained by the two models after supervised fine-tuning demonstrate this use within the scope of PinpointQA. These results should not be interpreted as evidence of broader improvements in general video understanding or unconstrained embodied reasoning.

E.2 Limitations

PinpointQA is constructed from ScanNet++ and ScanNet200 and is restricted to real indoor scenes and the selected vocabulary of small object categories. It therefore does not establish generalization to outdoor environments, other data sources, or unconstrained object vocabularies. Evaluating generalization across datasets and environments remains an important direction for future work.

The benchmark focuses on spatial grounding from video observations rather than general temporal understanding, open-world object tracking, or complete scene interpretation. Its four tasks end with natural language and structured spatial prediction and do not directly evaluate downstream navigation, manipulation, or physical interaction. PinpointQA should therefore be viewed as evaluating a perceptual and spatial capability that precedes these embodied behaviors rather than the behaviors themselves.

The QA pairs are instantiated from intermediate spatial representations using controlled construction rules. This design provides consistent and reproducible supervision across the four tasks, but it also produces more regular question and answer patterns than unconstrained human language. The benchmark consequently prioritizes reliable spatial grounding and comparability over linguistic diversity. Finally, the released benchmark does not redistribute the original videos or raw scene assets, so complete reproduction requires access to the corresponding source datasets.



Scene: 062e5a23a6

TPV

Question: In the entire scene, did the **mouse** appear?

Ground Truth: **Yes**

NRI

Question: Excluding the surface it is resting on, which of the following objects is nearest to the final location of the **mouse** in this scene?

- A. glasses case B. monitor
C. office chair D. laptop bag

Ground Truth: **C. office chair**

FSD

Question: Please describe in detail the final location of the **mouse** in the scene, including its distance and relation to nearby objects.

Ground Truth: The mouse is on the table. It is also to the side of and above the office chair, about 6 cm away. It is further near the glasses case, about 26 cm away.

SSP

Question: Please output the final location of the **mouse** in structured JSON format.

Ground Truth: {"target": "mouse", "support_surface": "table", "references": [{"object": "office chair", "relation": "side_above", "distance_cm": 5.8}, {"object": "glasses case", "relation": "near", "distance_cm": 25.6}]}

TPV(Cambrian-S-7B)

Question: In the entire scene, did the mouse appear?

Output: **No**

NRI (GPT-5.4)

Question: Excluding the surface it is resting on, which of the following objects is nearest to the final location of the mouse in this scene?

Output: **D. laptop bag**

FSD (LLaVA-OneVision-1.5-8B)

Question: Please describe in detail the final location of the mouse in the scene, including its distance and relation to nearby objects.

Output: The mouse is located on the desk, approximately 23 cm away from the edge of the desk. It is positioned near a black laptop, which is about 41 cm away from the mouse. The desk also has a white keyboard placed directly in front of the mouse, at a distance of about 15 cm.


SSP

(SenseNova-SI-1.3-InternVL3-8B)

Question: Please output the final location of the mouse in structured JSON format.

Output: \texttt{{{"target": "mouse", "support_surface": "chair", "references": [{"object": "chair", "relation": "on"}, {"object": "desk", "relation": "below"}]}}

Figure 1: Example 1. Errors across TPV, NRI, FSD, and SSP for target *mouse* in scene 062e5a23a6. The examples show an incorrect target presence decision, an incorrect nearest reference, unsupported references and distances in FSD, and inconsistent supporting surface, reference, relation, and distance fields in SSP. Green text agrees with the ground truth, while red text indicates incorrect or unsupported content.



Scene: 0b031f3119

FSD

Question: Please describe in detail the final location of the **mobile phone** in the scene, including its distance and relation to nearby objects.

Ground Truth: The mobile phone is above the tripod. It is also to the side of and above the table, about 23 cm away. It is further to the side of and above the keyboard, about 27 cm away.

SSP

Question: Please output the final location of the **mobile phone** in structured JSON format.

Ground Truth:

```
{
  "target": "mobile phone",
  "support_surface": "table",
  "references": [
    {
      "object": "tripod",
      "relation": "above",
      "distance_cm": 2.4
    },
    {
      "object": "table",
      "relation": "side_above",
      "distance_cm": 23.1
    }
  ]
}
```

FSD

Kimi K2.5

Output: The mobile phone rests on the white desk surface, positioned approximately 15 cm to the right of the black computer monitor and about 25 cm in front of the wall-mounted power outlets. The phone lies flat on the desk, separated from the keyboard by a small gap of roughly 20 cm.

InternVL3.5-8B-Instruct

Output: The mobile phone is on the table, 23 cm away from the keyboard and 41 cm from the mouse.

SSP


Kimi K2.5

```
{
  "target": "mobile phone",
  "support_surface": "desk",
  "references": [
    {
      "object": "keyboard",
      "relation": "next_to",
      "distance_cm": 15.0
    },
    {
      "object": "monitor",
      "relation": "near",
      "distance_cm": 25.0
    }
  ]
}
```

InternVL3.5-8B-Instruct

```
{
  "target": "mobile phone",
  "support_surface": "table",
  "references": [
    {
      "object": "keyboard",
      "relation": "side_below",
      "distance_cm": 15.0
    },
    {
      "object": "mouse",
      "relation": "side_below",
      "distance_cm": 12.0
    }
  ]
}
```

Figure 2: Example 2. FSD and SSP errors for target *mobile phone* in scene 0b031f3119. Kimi K2.5 and InternVL3.5-8B-Instruct produce readable or parsable predictions but replace the ground truth references, spatial relations, and distances with incorrectly grounded alternatives.



Scene: 281ba69af1

FSD

Question: Please describe in detail the final location of the **headphone case** in the scene, including its distance and relation to nearby objects.

Ground Truth: The headphone case is on the table. It is also to the side of and below the speaker, about 12 cm away. It is further to the side of and below the monitor, about 17 cm away.

SSP

Question: Please output the final location of the **headphone case** in structured JSON format.

Ground Truth: {"target": "headphone case", "support_surface": "table", "references": [{"object": "speaker", "relation": "side_below", "distance_cm": 12.0}, {"object": "monitor", "relation": "side_below", "distance_cm": 17.1}]}

FSD

GPT-5.4

Output: The headphone case ends up on the wooden desk along the wall, near the left side of the tabletop. It's right beside a folded gray cloth and about 20-30 cm from a white tissue packet, with a plant just off to the desk's right.

LLaVA-OneVision-1.5-8B

Output: The headphone case is located on the floor, approximately 23 cm away from the edge of the desk. It is positioned next to a black chair, which is about 41 cm away from the desk. The case appears to be resting on a blue towel that is spread out on the wooden floor.

SSP

GPT-5.4

Output: {"target": "headphone case", "support_surface": "desk", "references": [{"object": "tissue box", "relation": "next_to", "distance_cm": 10.0}, {"object": "door", "relation": "near", "distance_cm": 20.0}]}

SenseNova-SI-1.3-InternVL3-8B

Output: {"target": "headphone case", "support_surface": "chair", "references": [{"object": "chair", "relation": "on", "distance_cm": 0.5}, {"object": "table", "relation": "next_to", "distance_cm": 1.7}]}

Figure 3: Example 3. FSD and SSP errors for target *headphone case* in scene 281ba69af1. Although the target label is preserved, the predictions contain incorrect supporting surfaces, reference objects, spatial relations, or distances.

Mode

Annotation

Test

Test Set

Raw

Model

GT

Progress

11 / 20

Sample Info

Set: GT
ID: 20871b98f3
Images: 4

Timer

00:13

Session Summary

GT run ready. Results will be saved after the final sample.

QUESTION

Where is the scissors? Please click on its location.

GROUND TRUTH

The scissors is to the side of and above the dish drying rack, about 4 cm away. It is also to the side of and below the neon light, about 9 cm away. It is further to the side of and below the kitchen cabinet, about 12 cm away.

Image 3 / 4

Select a frame, click the target location, then save to continue.

← Previous



2.png



3.png

Next →



4.png

Start Test

Save Answer

Answer point selected on image 3. Save to continue.

Figure 4: Human evaluation interface in the GT-Assisted setting.